

An Interview with Eric Seufert About Models and Ads, and AI's Upside for Humanity

Thursday, May 28, 2026

[Listen to Podcast](#)

Listen to this post:



0:00 / 1:06:07

Good morning,

This week's Stratechery Interview is with [Eric Seufert](#). Seufert is, in my opinion, the world's pre-eminent digital advertising analyst, and [Stratechery's most frequent interview subject](#). Seufert writes full-time at [MobileDevMemo](#), and is well-worth the subscription. Seufert is also, as of today, the owner of a Master's degree in Applied Computation at Harvard.

Today's interview starts with an overview of Seufert's master's thesis, which described how companies can build models to measure the effectiveness of creative (thus unlocking the value of generative AI) using ad platform black boxes. We then move on to why chat checkout failed, how Google has shifted search to AI, and the rapid evolution of ChatGPT ads. We discuss Meta in two parts: first, why CEO Mark Zuckerberg is failing the company by not advocating for ads, and how the company is using foundational models to transform ad targeting. We also touch on Apple, Amazon, and AppLovin, and close with a fascinating thesis: is being pro-advertising correlated with being optimistic about AI and its impact on humanity?

As a reminder, all Stratechery content, including interviews, is available as a podcast; click the link at the top of this email to add Stratechery to your podcast

player.

On to the Interview:

An Interview with Eric Seufert About Models and Ads, and AI's Massive Upside

This interview is lightly edited for clarity.

Topics:

[DeCANT](#) | [Meta's Communication Problem](#) | [The Agent Checkout Flop](#) | [Google's Successful Gambit](#) | [ChatGPT Ads](#) | [Apple, Amazon, and AppLovin](#) | [The Prosperous Society](#)

DeCANT

Eric Seufert, [welcome back](#) to Stratechery.

Eric Seufert: Good to see you, Ben. Thanks for having me.

Well, this is a big congratulations, I believe about the time this is being posted, as people are listening to this, you are walking across the stage in Cambridge to get your Master's in Applied Computation at Harvard. Congratulations!

ES: Thank you very much, that's right, Thursday the 28th.

I'm not sure how much you've talked about what you've been doing exactly, but as part of your Master's, you just wrote a thesis which you just posted about on Mobile Dev Memo, [we'll have a link](#) — I think it's super interesting, I'm very excited to talk about this. So what you did is you developed something that you're calling DeCANT, a Deep Creative Attention-based Network for pre-Testing. Explain, because this actually touches on — it's funny, now that I've read at least the summary of your thesis, I go back through [your archive](#) and I'm like, "Oh, he's been pre-

writing bits about this for a while now”, because it called back to certain ones, particularly around generative AI and the testing problem and all those sorts of things. So this is your stage, you already defended it, you’ve had practice, now you have to defend it to the Stratechery audience, give us your thesis.

ES: Well, maybe a more discerning audience, that’s a little bit intimidating — thank you. Maybe I could just maybe describe some of the background, I haven’t talked about this at all, this is just something I’ve been doing in my personal life.

I know! It’s very exciting that we’re going to be the first outlet for this. I thought your summary article was great, it was one of those ones I immediately got what you were talking about right away, yet also felt very awed by the complexity of what it was going to go into it, so it was a good summary. So anyhow, I’ll stop interrupting, give me the lowdown.

ES: Three years ago, I think like a lot of people, I kind awoke to this idea that we’re standing at the precipice of this inflection point, not just in digital advertising, but I think in technology broadly and society broadly, and I got nervous that I just wouldn’t be able to operate at the frontier to provide commentary without better understanding machine learning, particularly in the digital advertising space. I just thought it was going to become more and more important to deeply understand these fundamentals, so I applied to go to graduate school and the idea kind of snowballed. I actually looked at like classes at a local community college and then I was like, “Well why don’t I just throw a Hail Mary here and apply to some more prestigious—”

(laughing) Went to school in Boston, as it were.

ES: Just outside of Boston and yeah, I got in and we ended up having to defer a year because my wife got pregnant and school would have started on the due date of my son, so we deferred a year, and then packed up moved across the country and it was just a really fantastic experience — but it was nice to go into it with some domain expertise too, to understand what I wanted to apply it to. I

think it made it a lot more concrete and a lot more actionable, but I don't think I would recommend this for anybody who is in my particular set of circumstances with remote work and my wife works remotely, and my wife is like a saint and allowing me to do this kind of stuff and young kids and all that, but it works for me.

So what I wrote about was we're sitting in a stage of the digital advertising market where essentially all of the big platforms are converging to this total end-to-end automation paradigm.

The black boxes.

ES: The black boxes. But if you kind of dive into a lot of the recommendation system research, which is essentially like the same as ad serving, it's all the retrieval and ranking methodologies that are applied for selecting ad candidates, it's the same as it's used for selecting content candidates in just normal recommendation systems.

It's always been weird, it happened for historical reasons, that these were two different tech trees, but it's been clear for actually a long time that if you're recommending an ad, you can recommend content and "What's the actual difference?"

ES: Right, it's just the objective function that's chosen. One is, "I want to optimize for this conversion", and then the other is, "I want to optimize for some sort of engagement metric", but essentially you're using the same technologies, the same methodologies, and so if you dig in all that, it's gotten supremely sophisticated.

But all of this sits behind this event horizon past which light cannot escape, with the light being your visibility into what the ad platform is doing. It's totally opaque, and it's gotten to the point where like you essentially have one lever of control, maybe two. The one lever of control is the creative and then the other

lever of control is whatever your objective is that you tell the ad platform. And I say “maybe one”, because you might not have creative.

That’s right, because generative AI’s coming along and it’s the most obvious application of just make the creative for you and it’s seemed like it’s been the case for a while, that if you’re an advertiser, “Just embrace the void”, has been the answer and that’s part of what makes this interesting. It’s kind of like a, “Maybe not, maybe that’s not the only way”.

ES: I think it is, that’s ultimately the point of this DeCANT’s methodology. But yeah, what are you going to do? Are you going to fight the ocean?

Right.

ES: This is the direction of travel, you are a single advertiser, you control some amount of budget — this is where things are heading. It’s obviously better for the platforms, it is probably better for you, so whether you like it or not, whether you would agree with it or not, whether you think things could be done in a superior way or not, it’s kind of irrelevant, this is where things are headed, and it’s obviously working out for the platforms, I’m sure we’ll talk about earnings in a second.

So if you accept that, and if you say, “Look, I have no control over this”, and also, “I understand that the cost of creative production is going to zero”, whether the creative production is handled by the ad platform itself, or whether I’m still creating my creatives, if I’m just using third-party tools to do that, it’s trending towards zero.

Now, there are a lot of reasons to believe that many advertisers will reject having the platforms make their creatives for them. For one reason, there’s brand conflicts. For another reason, you’re going to work with more than one platform and you want to have consistency across them. There are many advertisers that won’t just surrender everything to the platform, they’ll still keep control of the creative production.

So if you do that and if you say, “Look, I can’t control what happens beyond that event horizon, light can’t escape, I have no idea what’s happening behind the scenes, this is a total black box”, then you could just treat the platform essentially as this end-to-end system. Essentially, it’s a teacher model, and you could train a student model through behavioral distillation and just say, “When I submit my creative to you, and I apply it in this context” — the context being the country it’s being run in, which you can control for, the language parameters that you’ve set there, the date that you’re deploying the creative, and that’s kind of it, “If I have this context-conditioned creative, can I predict what the ROAS is going to be and distill off of all this information that I’ve gathered from the platform with my own historical performance?” — so that was the premise of the paper of the thesis. And it turns out that, what I empirically showed support for was that, well, this type of deep neural network can do that better than a tree-based method, which I used XGBoost as the best performing baseline, which is kind of like the canonical go-to when comparing performance of a deep network versus a more traditional ML method and so that’s what it showed. And I think there’s like a lot of cool applications for this, it could sit in between.

Just making sure I understand it, my summary of this, and I think you made a very important distinction — there’s one world where you just give Meta a price and they generate the ad, they do literally everything, and they deliver on the price and you know they’re getting better margins on it, but there’s nothing you can do about it because you’re still getting the best performance by doing that, that’s sort of like the default end state in one area.

But your bit about, “You want control of your creative, at the end of the day, your brand is your brand, that’s the most important thing”, but then number two, your point about you want to have consistency across platforms, which is this is where a lot of these sorts of total internalization efforts do fall down because at the end of the day, you’re the entity working across platforms. Meta doesn’t care about YouTube, YouTube

doesn't care about Snap, or whatever it might be. So you want to have control, but, and I think this was in your paper, so I'm going to add it to you, just doing it yourself, testing all your creative, costs a lot of money and is very wasteful, so that's the benefit of basically utilizing them as a teacher to train your own model. You can have a good idea before you run it with real dollars and sense what's going to perform and what's not, which lets you squeeze out a bit of that margin and efficiency that they're completely internalizing and maximize your return on the platform.

ES: Well, I think it's a little bit more than some incidental amount of money. A lot of advertisers use 10% as a rough benchmark for testing budget, now that's 10% with no ROAS accountability, so now imagine clawing a lot of that back because I'm pre-testing, I'm pre-testing with the model that's been distilled on their model, it's a lot.

10% is a lot if you're talking about these ads, yeah, that's a great point.

ES: But the thing is, think about from the platform's perspective, they might not like this methodology because, "Hey, wait a second", all that testing budget was like free money that they were being handed because it doesn't perform, but they don't care, this is just the cost of doing business, right? So that 10%, they wasn't accountable for ROAS, that got absorbed into just COGS or whatever line item bucket. But now, if they're actually looking at that as cost savings, well, that's going to be taken away from us. So I could see where the platforms might be a little bit hostile to this idea.

But is there anything they could do about it? That's sort of the brilliance of this model that I see which is it's embracing the black box nature of what the platforms are doing and, to your point, it's going to have to be retrained fairly frequently but you argue that's not going to cost very much. This is why I'm excited about it, it seems like something that's a pretty genuine innovation that is going to, if anything, get better over time. Am I wrong to be so optimistic?

ES: I don't want to oversell the model per se, but I think the application and the implementation is novel.

It's the concept though. It's using the network's ad server as a teacher for your own models, because they've reduced so many variables that actually it's just like an objective function.

ES: Exactly, and the idea of pre-testing, right? So just to test against some sort of model that should mimic what the broader system does.

Exactly.

ES: But I think there's another piece here too — if I'm talking about this to a room of brand marketers, I say, “Hey, good news, this allows you to take control of creative production and control your brand messaging everywhere”, they get really happy. If I talk to the hardcore DTC performance marketers, I say, “Good news, make everything, build everything, you can just shoot the fire hose at this thing and it's not going to blow up your testing account, it's not going to blow up your testing budget, but you should be making everything, you should be exploring the entire surface area of all possible messaging and pixel combinations that could represent your product — shoot the fire hose, blast everything at it because the cost of producing creative is essentially zero”. Now, you don't want to hand that over to Meta or to whoever and I don't want to single out Meta because it's Meta, Reddit, Google, Pinterest, TikTok.

Meta's Communication Problem

Right. Meta's the stand-in for everyone because they're also the best at it, so we get to use them.

ES: It's because of [Meta CEO Mark] Zuckerberg's [quote on your podcast](#).

Which was?

ES: “Just give us your credit card, we'll take care of creative, we'll take care of measurement, we'll take care of everything else”.

That's right. Yeah, that's been the bull thesis for a long time.

ES: It's been the bull thesis when you talk to the DTC guys, it's not the bull thesis when you talk to the brand guys, and I think that's part of the issue here.

Right.

ES: It's like Mark Zuckerberg, I respect him very deeply as an executive, especially in terms of someone who just accomplishes things very ruthlessly, he's not a great communicator.

Nope, about ads in particular.

ES: Exactly. And he does a really bad job, I think, of communicating the vision for the ads platform, if he even deigns to, because a lot of times he just blows that off.

I had a whole thing about this, this has been one of my theses for a while, is the entire industry has been hurt and Meta in particular has been hurt by the fact Mark does not go to bat for ads. I'm a much bigger public advocate for Meta's ads, which I think are a societal good, it makes entirely new industries possible, new niches, make customers aware of them. In a way that I see search being more of a skim that's a little icky and I'm like, "Why am I out here making this case far more vociferously and effectively than you are who likes to pretend that ads don't exist?", it drives me crazy!

ES: Let me just pitch a bit of a thought experiment to the audience, I don't want to put you on the spot, so this is for the audience not for you, although I think you probably know the answer. What's the name of Meta's COO?

Boz [Andrew Bosworth], right? Oh he's CTO...who is the COO? That's a good question, I don't even know.

ES: Javier Olivan.

That's right, I did know that, but yes I did not know it off the top of my head.

ES: But you knew Boz, you know Zuck, you probably know [Meta CFO] Susan Li, but who is the last COO you knew about? It was Sheryl [Sandberg]. You heard her on every earnings call — that's the issue. Sheryl was the champion for the advertiser, she talked about it on every single earnings call, 10 million advertisers.

She'd always come in with her five anecdotes of advertisers that are doing awesome, it's a great point. Google does that, they always have like — and you kind of roll your eyes as you go through it, but there's a repetition to that of, "Look at this real-world result that is positive", that is totally missing, it's a great observation.

ES: So just to go back to the point about the platforms, I think also the platforms would be happy if the advertisers are happy and they're making more money and if they're deploying better creatives that scale better and they're not actually — there's not a lot of dead weight loss on creatives that don't work then, they're probably going to be spending more money and so probably this is win-win for everybody, just this kind of approach.

And also if the advertisers are spending all the money to generate the creative, saving their compute resources, I guess at a minimum.

ES: Right, yeah.

Number one, I think it's great paper, number two, I'm very interested to see what the implications of this are because I think it's very compelling, it applies broadly.

The Agent Checkout Flop

But there is a second big reason to have you on the podcast, I need to pull up the wheelbarrow to give you a massive heap of Being Right Points. From day one, you expressed strong skepticism about ChatGPT's checkout

product, I think we argued about it a bit the last time you were on. Fast forward, I'm not sure if that product even exists anymore, or at least it's very deprioritized, Walmart expressed a lot of dissatisfaction with their participation, which I took Walmart's participation as a positive. You were absolutely, completely, 100% right; where do you want to start your victory lap?

ES: I wrote a piece that got me more hostile feedback than any other piece I've ever written — including some of the more critical ones I wrote about Apple, which attracted a lot of negative feedback — called Agentic Commerce Is a Mirage.

Look, that was a clickbaity title, that wasn't the point I was making in the piece, the point I was making in the piece is that third-party independent agentic commerce is just not viable. It doesn't align with the interests of the platforms, it's probably not the best way for those companies, the chatbots, to make money. Advertising is superior and that it just, maybe it was an entry point to bootstrap data, but it wasn't going to go anywhere, that's not how people want to buy things, they need to have the platforms, it's not how they want to sell them, they're not going to cooperate if it's going to cost them ad revenue, it's not going to be conducive to doing all the RecSys [Recommender System] stuff that gets people to expand their cart when they check out and for all those reasons, Instant Checkout didn't work.

Right. And to your point, this isn't just like agents buying things, this is also like buying things from within a chatbot experience.

ES: So the fully automated agent buying on your behalf, I think, is non-viable and I just think that's not going to work ever. But I think the third-party independent agent of commerce where they're just recommending stuff for you and then you click through and they get an affiliate fee or actually they facilitate the checkout in the chatbot and they get an affiliate fee — I just thought, first of all, that's just

suboptimal relative to advertising, you get more economics in advertising and also it's more win-win if you use advertising versus the platforms are obviously—

It's very adversarial, yeah.

ES: Amazon held out, you needed Amazon's participation if this was going to go anywhere, they haven't, they still haven't, they've been excluding more chatbots if you look at their robots.txt.

But I do think agentic commerce, when it's bolted into the platform experience, is a viable solution and I think it is working really well. Rufus is working really well, Walmart's agent commerce tool Sparky is working really well, and so I think there's, I wouldn't even say that that's a future state, I think it's happening now, I think it's had success now.

And so the issue isn't necessarily the interface, it's the fact of these conflicting incentives where if you're on Amazon and using Rufus, Amazon still gets all the data benefits, the advertising surface area, all the sorts of things and everything's up to date. There's no integration issues, the cart is consistent, all these sorts of things.

The big issue was, yes, ChatGPT can surface you a product you didn't know about, but every other part of that experience — I think this is like, the more I thought about why you were right, there's always those stories that you hear about like the impact of millisecond delays in conversions. If you think about that, if that's a talking point on a site where you're already there to buy, almost definitionally any sort of integration, just the fall, this is what Walmart was talking about, they're like, "Our conversions through this are terrible, it's better to get fewer people clicking through to our site where we could actually make sure they finish the purchase", which I think is like is the perfect articulation of the point you were making.

ES: Right, and the AOV [Average Order Value] was lower, people were buying one thing at a time. You go to Walmart, you put a thing in the cart, you get a

recommendation for five other things you put some more of those in the cart, the average order value goes up.

And the other thing though, I think the bigger more fundamental thing was that the economics were just suboptimal to advertising, the economics just didn't work as well as advertising.

What do you mean? What do you mean like an affiliate fee versus versus an ad?

ES: The affiliate fee versus ads, because the ads you submit your true value as a bid, there's just more value to be had there. If I'm just getting an affiliate fee that's like an arbitrarily chosen percentage off of the top line and then ultimately you're going to end up preferencing this low cost stuff and so that could degrade the user experience, but also you're just going to make less money.

Having a bid allows everybody to say, "This is what it's worth to me to get that impression", and so you're going to get the higher value stuff — it's modified by the conversion rate, but you're going to usually tend to get more higher value things that get selected for versus just the affiliate fee, which is a fixed percentage is going to tend to have lower valued stuff because it's a higher conversion rate.

You've doubled down on your argument that AI advertising — this is somewhat related — should not look like search advertising. You explicitly wrote last month, Search Is the Wrong Mental Model for Chatbot Advertising, and I've largely agreed with you and made the case myself. One area, though, where I've gotten pushback from people is the questions about the mental state of consumers. It's like, to what extent can AI actually be like an Instagram experience where what's actually the difference at the end of the day between the content and the ads? People are on AI to do work, to be productive, display advertising is not going to work there. What's your pushback?

ES: I think anywhere you're applying attention is a good candidate for display advertising.

Right, and there's a vast umbrella and conversion capability underneath Instagram that is still very profitable.

ES: And also I would point you to mobile gaming, mobile gaming is very much a lean in activity, it's not a passive lean back consume Instagram videos activity, and you know ads work pretty well in mobile games.

Great analogy, yep.

ES: So does a lot of streaming too, maybe less so now, I think a lot of these streaming shows are designed for lean back passivity, but I just watched the new Jack Ryan movie, I could not make heads or tails of that narrative, it was really poorly written, and they did the thing that Matt Damon was talking about where they restated the plot in the dialogue like six times.

(laughing) And it's not just that you weren't paying attention, it's also that the plot didn't make any sense, so maybe they were just making it worse.

ES: It was just a nonsensical plot, but that's an aside.

So my sense is like, if an ad break is an ad break, and they do capture your attention, that's actually one of the arguments that AppLovin makes with its advertising product, which is why they think it's more attractive than social, because it's an actual ad break, you get a pause in the gameplay, and you're forced to watch some minimum number of seconds before you can move on, that's a really compelling idea. And again, if you tie this, if you anchor the ad to the context of the conversation, there's going to be some very large percentage of conversations that are not monetizable at all.

Right. Which is the case with search, by the way.

ES: Exactly. They said 80%, historically we're not monetizable at all and so you get so much more monetizability if you're actually just tethering this to whatever the person is doing elsewhere on the Internet.

Google's Successful Gambit

Has Google actually — is this an argument they've been dramatically under-monetizing Search? Because all these non-monetizable search results should have Facebook-style display advertising next to them, because why not?

ES: Yeah, maybe. I think what happened with Search — I mean first of all, search is fantastically profitable, it's probably like one of the best business model innovations of all time, and it grew 19% last quarter on a huge base, so I'm not trying to denigrate the performance of Search.

But I think Google Search may have been a victim of the MAYA principle, like an industrial design, Most Advanced Yet Accessible. And so if you change it too much, then it's not Google Search anymore and people don't recognize it and they know how to use it, so you could change it incrementally and that's what they did through what I've called **Google's Gambit** and this like pretty protracted evolution of that experience.

I'll let you finish, but I had Google's Gambit further down the interview, I'd love for you to walk me through that as well.

ES: The premise of **Google's Gambit** was essentially that they were taking this big bet — keep in mind I started writing this series like 18 months ago, so when they introduce AI Overviews — they're taking this big bet, "If consumers are getting accustomed to a chat-like interface, we have to adapt to that, but if we adapt too fast, then people will say this is a totally new experience, I might as well just go to ChatGPT" — this is disorienting, you can't just overnight flick the switch and have Google as Google AI mode, so you've got to take this in incremental steps.

The point I was making when I started the series was they're going to have to perform this Ship of Theseus transformation of Google Search into a conversational chatbot-like experience over time, because that's what consumers are getting accustomed to with chatbots. They can't do it all at once, and when they do that, though, they're going to cut off the open web, they're going to capture way more of the value, and they're going to turn Search from a distribution mechanic where they just distributed clicks and distributed traffic, into an engagement sync where people spend time in search.

That's a foreign concept 18 months ago, people don't spend time in Search, they go there for a specific purpose, they go there to go away, they go there to click away, and if they don't click away in the first link, that's a failure, that means Search did a poor job, so they go there to get the result that they want to click away to very quickly, and they needed to transform Search into something where I go there and I have a conversation. I gather all the information that it would have gotten downstream from some third-party destination, I gather that in the Search experience, but they can't, again, they can't flip the switch.

So what they did was they introduced AI Overviews, which displaced a lot of the links with just the summary. Then they introduced AI Mode, now they've finalized this process and they've made AI adaptability a core component of the Search experience, the direct Search experience, so this is the total Ship of Theseus transformation. Now it's happened and it's done and it's worked beautifully for them, they said they're monetizing at parity, queries are up, it didn't kill Google Search, queries are up, and Search revenue grew 19% last quarter.

So can I reframe my question then? It's not that Google missed a display advertising opportunity before, it's that now actually it's not just that they can do search advertising, but now they have a canvas and they have engagement in order to add on this new means of monetization.

ES: Yeah, exactly. And they did it with a catalyst of consumers becoming accustomed to the chatbot experience. I think absent that catalyst, it probably wouldn't have made sense and it probably would have been disorienting, but this was an established consumer behavior, established interaction mode, that consumers had been introduced to and so they had a lot more freedom, I think, to adopt it.

Because the other thing is like, you've been seeing the headlines about like Meta **is about to eclipse Google** in ad revenue, ex-tax, and if you back the TAC [Traffic Acquisition Cost] out, that Meta is closing in on Google overall ad revenue. I think it's overstated that Meta is going to eclipse a Google on this point, but when I talk to people about this, people say like, "Google can't risk it, look at Google Search revenues, Google Search revenues are so massive and they're so strong", I'm like, "Yeah, but they have this massive distribution benefit that they buy from the TAC".

But the thing is, so are Meta's — if you look at this idea that display is a tiny fraction of the overall digital advertising ecosystem is just a bizarre misconception, it's a massive part of the digital economy. So the idea that search is better, it's de facto better than display is silly. Keyword-based search, I think is probably going away, I don't think that makes sense in this sort of modern multimodal environment and so having the ability to de-tether just gives you weight from the actual context of the search, the conversation, gives you so much more freedom to monetize everything, to use the whole buffalo.

ChatGPT Ads

To take it full circle, this is why you're optimistic about the chatbot advertising generally, just because it works. There's just a certain aspect that people just don't believe in it, but they should.

ES: Demonstrably works, OpenAI had **\$100 million in ad revenue** in like three weeks after they launched the beta.

With terrible ads that we all complain about. Well, you're the one who's like, "Well, those are experimental budgets, let's see how that's real" — is it real?

ES: Yeah, it's real. I mean of course that was experimental budget, but they're targeting a billion this year and this is more of an actual modern ads platform if you look at how they're executing. They are executing blazingly fast, they are building this stuff really quickly.

Well, let's talk about it. So they just launched a CAPI, a pixel, and self-serve. Give us a quick overview of each of these and what it says to your point about how fast they're moving and where they're going.

ES: People pushed back when I wrote that article, but I was just reporting what they were publishing, which is, "Well the CAPI and the pixel had existed before", yeah they were in beta, they were in testing now, they exist for the self-serve

What is a CAPI [Conversion API] and a pixel just for our audience?

ES: These are both just measurement methodology.

So the pixel is what gets integrated in your website, it fires on load, it's a blank pixel, it's a 1x1 blank pixel that, when you land it fires, but actually the pixel really is that the power of the pixel is the combination of the cookie and the pixel, because the pixel fires, it tells you where you're at in the shopping journey and which product pages you're at. The cookies, what captures the UTM when you click through to say, "Okay, they came from this campaign ID, they have this click ID", and so it's a combination of those two things. It's really powerful that the pixel just allows that to be sent, the signal to be transmitted that you're landing at that page, because that's where the file gets called from. So there's measurement methodologies to know that a user hit this page, the user clicked on this, you've got the click ID and the cookie and then you've got the location of where the user's at, through their journey, so you can measure to see what they did after clicking.

CAPI is the server-to-server answer to that, and usually you'd introduce both, but the reality is I think people rely more on the CAPI now. CAPI was introduced by Shopify after Facebook and Apple went a couple rounds of whack-a-mole with ITP — ITP is Intelligent Tracking Prevention that Apple introduced into Safari. Keep in mind, Apple has blocked third-party cookies in Safari since 2017 and so what Facebook kept doing was finding little loopholes to get around that and so I think ultimately, Apple blocked first-party cookies from platforms that have any sort of link decorations, which would be used for tracking. They blocked those after 24 hours, like 24 hour expiration, which is too short to use for measurement. So what Facebook did is, “Ok, let's just get out of the browser, let me just set up a server endpoint, you tell me what the user is doing, and then I'll do some probabilistic matching on my back end”, which is most likely done via IP address and some other signals.

So that's more reliable, but you kind of need both because a lot of advertisers still, even to this day, there's a lot of tools to help with this, but a lot of advertisers still really aren't ready to set up the CAPI.

Pixel is like old school advertising, CAPI is like new school advertising — pre-Apple crackdown and post-Apple crackdown, you still need both.

ES: Exactly. The thing is, though, like what OpenAI did with the CAPI, this is just allowing advertisers to send post-click information about what the user did, so for measurement, but the thing is usually what you do with this and what a modern ad platform does, but what OpenAI doesn't do yet, is you optimize for those events, so you optimize the ad serving for those specific events that you're capturing. OpenAI doesn't offer that, they just created the ability to capture that information for measurement. So you could say, “Ok, we're still doing CPC [Cost-Per-Click] bidding, but we can tell you when you spent this much money, this many events happened, and then it's up to you to decide if that was profitable or not.

And that's just a matter of time, right? Just because they haven't built it yet.

ES: It's just a stopgap, but they are executing very quickly, so my sense is that update is not too far off, the actual conversion optimization for those events.

So by-and-large you've actually been heartened by the speed with which they're implementing these changes, even though their first launch maybe it really was, "We just have to get something over the fence to force us, it's going to be bad it's going to be embarrassing but that will give us the motivation to actually ship a useful ad product faster than we would otherwise"?

ES: Yeah, I wrote this piece when they first announced their ads offering which was meager and totally unimpressive, I wrote the piece called OpenAI should know better and my point was, yeah, they probably do know better, give them a little bit more than like 24 hours.

No I think in retrospect it was they just had to launch something and they needed the embarrassment function to force them to fix — it was such a bad product to start with but it's gotten much better already.

ES: I almost wonder if they just needed to get something out there to prove that the demand existed internally to get more support for advertising. Like, "Look, if we push this out, we've got \$100 million of commitments for this totally primitive system, think about what we could do with actual engineering resources". I think that might've been part of it, but I don't know that to be true.

Overall though, are you pretty optimistic about ChatGPT and the advertising opportunity?

ES: I would say I'm optimistic about the advertising opportunity if they keep executing as capably as they have. I agree with you, they should have done this a year earlier or more, I worry a little bit that Google set the tone for what AI or

chatbot ads should look like and they're doing a really good job and they've got just more interconnectivity. **They've got UCP now**, they've got all these interesting ways of essentially piggybacking off the existing ads relationships that they have to optimize the chatbot ads experience.

So if they set the tone there with a product that outperforms, and then again, to your point, why maybe it would have been good for OpenAI to have the first mover advantage here to just to set the stage from a consumer expectation point, a standpoint for what these ads should look like, but also like an advertiser expectation standpoint. Now that you've allowed Google to do it, obviously that is going to perform better because they're piggybacking off these existing relationships and this ties into the entire ecosystem. Again, it's not just AI Mode where these ads are being shown, it's everywhere. This is interconnected to the entire ad account, which is pervasive across the Google ecosystem, that's Google's real advantage here. Not only do you have the data from across this entire suite of products, but you have ad placements from across the entire suite of products. So there's a lot more information, but a lot more ability to deploy the advertiser's budget.

Now, if I'm saying, "I want ads in chatbot-like interfaces, and I can get that from Google", or, "I can get that from OpenAI", and each one of those requires its own integration, its own sort of data modeling, its own creative, but also with Google, I get all this extra for free, then you might just go with Google if you had to pick one.

You can already see in the sort of the attempt to develop these various protocols for ads and products, OpenAI kind of got too far over their skis, in part because they were headed in the wrong direction with Checkout and all these bits and pieces and then Google shows up a couple months later with Shopify and it looks like that's the protocol is going to win. Stripe switched over, they now technically support both and, yeah, OpenAI, by virtue of being late, they didn't get to define the playing field, they are playing on Google's playing field.

ES: You don't want to compare yourself to Google, you don't want to compare yourself to Google on advertising. You just don't, you don't want to, you don't want to force an advertiser to do that.

Meta's Models

Where's Meta in this? Are you encouraged by Spark? Are they going to be competitive here? Or is the actual Facebook opportunity, to tie back to what we talked about at the beginning, actually the product and the ads are all the same thing, every pixel is a monetization opportunity, stick with that.

ES: Products and ads are exactly the same thing, right? These systems are unified — they're not literally unified, but they are informed by the same research. And my sense is, I was just writing this today, but it's something I've been thinking about a lot lately, there's this quiet revolution going on right now within RecSys.

It's transitioning from this traditional DLRM [Deep Learning Recommendation Model] model, which was introduced by YouTube in 2016, so this predates the Attention Is All You Need. This idea of what's called the two towers, you've got a user tower and you've got a content tower, and then what you're doing is you're building these large embedding tables and you've got the user history and you've got all the possible content they could consume and you're doing just dot product similarity and you're trying to figure out what to show the user next.

That was tremendously powerful idea when it was introduced by YouTube in 2016, but this is giving way now to generative recommendation, and so this is also essentially invented by Google, I would say. They introduced this idea of Semantic IDs in 2023, this paper, this very famous TIGER [Transformer Index for Generative Recommenders] paper, and the idea here is that, I won't go too deep into it because it'd be a distraction, but instead of saying, "Here's all the possible content, and here's all the sequence of activities we have about this user, and let's condition the combination of both for every single possible piece of content that this user could consume", what if we just predicted the next

thing that the user would be likely to engage with? And what they do is they flip it around.

Totally flipping it around, yeah.

ES: So what they do is you can use an LLM for that, what the TIGER paper did was they trained a sequence-to-sequence model, but they trained an encoder and a decoder, and they did this codebook model with a methodology called RQ-VAE, [Residual Quantized Variational Autoencoder]. But what they did is they say, let's say you have a billion items in your content catalog, well, if you wanted to use like sequence modeling for that, or if you wanted to use an LLM for that, you need to create a billion, essentially, a billion representations of tokens. But you can't put that into a tokenizer because that would overload the vocabulary, you can't do a softmax on that.

And so what they do though, is they use this, the Google paper uses RQ-VAE, and what they do is they train code books, and so what they do is they use quantization to basically — what happens is you create these hierarchical code books where you, let's say it's a code book of size 256, so you've got four, so you've got 1,028 possible options, but it's the combinations that create the variation volume, and so you train these codebooks. Codebook one, I'm just basically grouping things, I'm classifying things into 256 classes. The second one, I take the residuals on the first one and I classify the residuals and I take the residuals of those, I do the third one and the fourth one. Right now, with a combination of four codebooks of 256 size each, I can capture the entire size of my catalog. Now, each of those turns into a token, so instead of 10 million or whatever tokens that I have to push into the tokenizer, I've got 1,028, and I push those into the tokenizer, and then I mid-train an LLM to accept those as potential output as next token prediction. So I say, "Here's a sequence of the things that the user engaged with, predict the next one as a token".

That has been revolutionary, I think. And now you're seeing a lot of these big advertising platforms publish papers are saying like, "We're going to rebuild our

entire architecture around this”, and it’s not just Google, it’s Meta. They’ve talked about LLM for ad ranking, they talk about it in their earnings, that’s the altitude that this has reached. They talk about the power of LLM for ad ranking, but also it’s Kuaishou in China. And so the thing is, if you think about all of the semantic and the world knowledge, what Meta says when they talk about this in earnings, all the world knowledge you bring to bear on predicting the next thing the person will be interacting with if you use this methodology, so I think there’s like a big leg of growth that is yet to even really be tapped into with the application of LLMs here.

But also think about the value that owning the LLM model has in that world, because I think we’re probably entering a world where, hey, maybe the most state-of-the-art LLMs, and we’ve already seen this with Mythos, “The most state-of-the-art LLMs, we’re going to keep those for ourselves, we’re not going to make that available”.

Mythos and Meta Spark came out the same week, which on one hand would seem negative for Meta, their model is relatively far behind — in my head, I’m like, this is the greatest endorsement of Meta building their own models that you could ever get, it’s existential. They absolutely have to be doing it, and the fact that Spark is not terrible means they’re on the right path, so everyone should be happy about this.

ES: Yeah, it was competitive. It’s not the tip of the frontier but it was competitive.

That’s fine, it’s on the right path, and it’s a path they have to be on.

ES: They own it, you can’t block them from using it, they own it. And what happens if owning that means that you are the best ad platform because you have the access to that LLM to do this kind of mid-training? I think that people discount that, but again, it’s because Zuck doesn’t talk about this, he doesn’t talk about why this is important, all these investments that they’re making in **GEM** **[Generative Ads Model]**. GEM is a foundation model, I think people think like GEMs a tool, GEM is a foundation model.

Has phenomenal results, that in and of itself justifies all the infrastructure investments they're making.

ES: But if you think about, “What does a research team do?” — they make foundation models! Well, they made one, they made GEM. And the thing is, if you look at the application of GEM, you can't use GEM for inferences, it's too big. Meta published a paper called SUM [Scaling User Modeling] in 2024. But they talked about all the places where these point solutions are utilized, all the different ranking scenarios, it's like 40. There's like all these different models that are distilled from GEM and get applied at different parts of the Facebook and Instagram experiences where ranking is induced, so there's a lot of applicability here.

And the thing is, it spans both of the products in every single surface area. I think Zuck should be talking about it more, and I think when you talk about Superintelligence, you're just raising the specter of the Metaverse. You're making people think, “Oh come on, this is a cosmetic exercise, this is just Zuck tilting at windmills”.

This is what ultimately, to your point, ads and recommendations have always been the same problem, but this is like the pivot point to actually make them into the same solution where you flip it around and you're understanding the user and what they actually want to see next. And then it's like inevitable that it's tied together, you're almost forced into the optimal organization.

ES: Another piece here, I was reading some analysis yesterday or the day before, and someone was making the point that, I genuinely don't remember who was writing this, but someone was making the point that it's all the adoption, all the economics of these models are being absorbed at the frontier, so it's the latest versions of the models that are capturing all the value and doesn't that show support for the idea that this will continue to move in that direction.

My sense is that's probably true because a lot of people are adopting these models and they're integrating them into their own commercial workflows for the first time, but what happens when they're not? If it's already done and there's a new frontier model, would you upgrade? And I **think you made the point** on Flash 3.5, it uses more tokens and it's a higher cost per token why would you upgrade? If I built a tool that uses Flash 3 and it works perfectly, I tested it, it gives me the output I want, what incentive do I have to upgrade in the first place, especially if it costs more money? And then the other thing I think people don't recognize here, it's like every time I query from a model, it's essentially like sampling from a distribution. Well, if I'm querying from a new model, it's querying, it's sampling from a different distribution, I'm going to get different results and so I might have to do a lot more than just change out the variable name. I have tools that I've built, the model that's used is a variable name, I could just change that and I'd be querying from 3.5, but that's going to change the entire output, I'm going to have to then change the system prompt that I use, I'm going to have to test it a bunch to make sure I'm getting what I want.

If I've already done that and I'm getting stable results from 3, then I don't know why I need to upgrade. Now, think about like, what happens if these models are actually so powerful for my use case, which is RecSys, which is ad ranking, it's re-ranking, that I don't want to make them available to anybody else, this is proprietary to me, I spent a lot of money, but I'm getting a lot more in value out of it in my 5%, 6% every quarter improvements in time spent — that's RecSys, in conversion rate, that's ad ranking. If I'm getting these results, why would I want to give that away for all of these competitors to then just do mid-tuning or fine-tuning or distillation?

Apple, Amazon, and AppLovin

We're on the same boat, Meta has to be doing their own models. To your point, if you're happy on Gemini 3 Flash and then Google pulls it, then you're screwed, go to the 3.5 whether you want to or not. Does this impact at all your view of Apple and them not doing a foundational model? Or is it

because yes they have advertising, but it's not the existential aspect of their business like it is for Meta that they'll be okay?

ES: Apple engages in ad-washing, they outsource all their advertising, Apple is highly reliant on advertising as a business, I wrote a piece about this, if you examine the contribution that Google's rev share makes plus their own advertising endeavors—

Ad-washing is a great way to put it, they're making a lot of money from ads, it's just flowing through search partnership deals.

ES: It's like 100% margin, they got the one guy that they pay to sign the papers every year, that's it, and he just sits in office for 364 days.

Here's what I think is happening with Apple — first of all I do think it's a misconception that Apple does not build these models themselves, they have a three billion parameter LLM that's integrated into the iPhone, it's on device that's their own model. You can use it through the Core ML framework as a developer, you can access that, you can access vision transformers, you can access BERT, they have all of this embedded on the phone.

My sense is Apple has to subvert its own privacy messaging if it starts building frontier models because where do you get this data from? You collect it. If you want actual unique data, you collect it from your users. Are they really going to say, "Privacy, that's Apple, except for we capture every keystroke and we use that to train our own models". Plus, they've kind of backed themselves into a corner with their application of DP everywhere, Differential Privacy. Even if you put that aside, I think this would just subvert its own privacy messaging, but I think also, you can do the LLM-washing just the same as you've done the ad-washing and have pay Google Gemini now, but at some point, if people get familiar enough with all these different providers, they're going to have a model garden type thing where you, instead of selecting your default browser, you select your default LLM when you've got subscriptions to, and that's going to be highly

profitable positioning, highly profitable distribution. So I think you're going to get the model providers paying in the same exact way that Google pays now.

But there's another piece to this, which I think also, the reason you couldn't put a frontier LLM on the phone is really the memory constraints, the KV Cache would be too big, you wouldn't be able to process that, but at some point, I think you'd probably catch up. Maybe not like a 1.5 trillion parameter model, but it could be pretty large.

This is the great thing about a bubble is you're getting innovation and invention at every part of the problem and people are working on the algorithms, they're working on the KV Cache, they're working on the memory, they're working on the quantization. All these bits and pieces, there will be breakthroughs, even if we don't know what they are now.

ES: Yeah, TurboQuant's like what, four months old, three months old? This is brand new, all this stuff is happening at lightning speed. I think they're probably best positioned to just wait and pay Google the pittance of a billion a year nominal fee to get arguably the best performing model integrated into Siri and Apple Intelligence. And then at some point they could have some mix of on device versus maybe their own model.

To go back to OpenAI and advertising and the agent bit, we both noticed [Amazon CEO] Andy Jassy talking about agentic advertising and third parties, it was kind of weird on the last earnings call. Amazon and OpenAI signed a big deal, they're doing this whole platform for stateful agentic work in addition to OpenAI now being broadly available on Amazon. Do you think as part of that deal, is there some sort of come-to-understanding bit about advertising and ChatGPT?

ES: I do. That deal was rumored for a long time but before it was public, I wrote that they're probably are going to do the investment and I think what that means is that Amazon will be able to serve the ads on ChatGPT. OpenAI did all these deals, they did one with Criteo, a number of other providers, where they bid in.

It's not a traditional RTB [Real-Time Bidding] partnership where they're buying impressions and fulfilling them, it's more they're just bidding in and then the ranking algorithm that OpenAI's built does the rest. But my sense is what you do is you get Amazon structuring something that looks similar to what they announced with Snap and Meta, but never really went anywhere.

Right. Where you sign in with your Amazon account so it's linked together.

ES: The biggest issue with Amazon is that they just never got social to work. They never got a standalone, beyond the e-com platform, they never got like a standalone engagement sync — I mean Twitch is something. They've always been searching for that and in my sense — and they've done deals with TikTok — my sense is they do is they are going to have a similar sort of product alliance with ChatGPT that allows them to just fill out the ads there.

You talk about Amazon's advertising advantages and I think you were talking about in the context of like video, about all the people that they have signed in, is this something that they'll actually extend this out for real?

ES: Yeah, it's a massive advantage, they have this massive identity spine of almost the entirety of the United States. If you look, they said we have coverage of 90% of the United States, they said that at their Upfront.

And they're all signed in, it's amazing.

ES: Signed in and they've got the most — this is not like, "Which cat videos are you watching?", this is, "What are you buying?", this is retail data. This is the most incredible, valuable data set ever accumulated. So yeah they could just, they don't necessarily need to own these products where they serve the ads, they just need access to the inventory. Honestly, there's probably very few companies that are better positioned to serve it and so it's kind of win-win.

This solves the issue of, “What sort of advertising works in ChatGPT?” — what sort of advertising works is products, “This is the product you’re looking for”, and, “Here, you can go buy it”, and if you just offload that entirely to Amazon, that is very compelling.

ES: And probably better than you’d be able to accomplish on your own. Now, keep in mind, Amazon’s also got like a pretty scaled DSP [demand-side platform]. They’re doing this on the open web, and so they’re doing this everywhere. It seems like the advertising strategy, the growth strategy, is not trying to get more surface area, it’s trying to serve more surface area.

Do they end up as the — we were talking about Meta passing Google, is Amazon going to be the biggest in the long run?

ES: Maybe in the very long run, they’re quite a ways behind and they’re pretty distant third, but yeah, they’re growing at an impressive clip.

Speaking of finding more inventory and surface area, is AppLovin launching a social media network?

ES: They did, it’s live, you can go download it, use my referral code.

Here’s our usual AppLovin discussion, this company continues to kill it. What is working for them and why would they go in this direction?

ES: There’s a lot of value in owning more inventory, so having more supply would be beneficial to them, it’s called Gist.

It’s like a Pinterest sort of thing.

ES: Yeah, it’s like a Pinterest meets that Chinese app—

RedNote.

ES: RedNote, yeah, it’s kind of like a hybrid of those two, I’ve used it. It wasn’t super compelling, but we’ll see where it goes, it’s still early days. You have to have

a referral code to even get into it so it's not meant for broad-based use.

There's a couple of things — they used to own a pretty scaled gaming catalog, they sold that to a company called Tripledot and the question was, “Well, if they've offloaded their own inventory, their owned content, why would they build this?” — there's a big difference. Gaming, it's just like this constant flywheel of having to buy attention. You buy a user and you get that user, you get a 90-day LTV, 180-day LTV, but with the social network, you get a probably much more viral distribution, so you get more than just one user and a sticky social network person could use it for years and years. So they are going to do a lot, I think they're going to do a lot, of user acquisition if this proves out, this could just be a total trial balloon.

Right. It sounds like it's an experiment, and if it works, then they're going to become big ad buyers themselves, that's been shown to work.

ES: Just having a social surface too, probably will attract other types of advertisers that aren't advertising, maybe that fills out their e-com initiative — this is just an experiment.

I think there's a lot of value in owning a scaled social media platform. Again, they bid on TikTok, but I think this is just an experiment, I think if you look at AppLovin's growth, it's still gaming, it's still almost all gaming, and now their e-com initiative is growing pretty quickly, and it's not in GA yet, that's coming very soon. It's still just gaming and the whole gaming category is seeing growth now where it didn't as a result of ATT for a number of years, but it is now. But also, they've made a lot of investments into the infrastructure that have worked out really well.

Where else can they go? Building your own social network — good luck to you, I applaud the effort, seems like a long shot, but is that actually a bearish indicator that they're even thinking about something like this? You mentioned they have this e-commerce business, it's doing better than I

think a lot of people expected, but they do still feel stuck in gaming land. Is there a way out of there?

ES: CTV [Connected TV]. They bought a company called Wurl a couple of years ago. That was a CTV SSP [Supply-Side Platform]. CTV is a pretty massive market and so I think that they have all the bones to expand there. But I think once you get that nexus of algorithmic precision and efficiency, you can apply that to a number of different places.

They just need to find all these areas where it's not necessarily logged in users, you have to figure out who they are, what they're interested in and they're so much better at that than everybody else, they can bring that to bear.

ES: Exactly, and I think CTV is probably another area that they explore.

That makes sense.

ES: It's kind of up for grabs, I think.

The Prosperous Society

Well, you mentioned you've been thinking about the big picture, zooming out. In addition to all, we got very tactical with a lot of these questions, you put together a four-part podcast series expressing AI optimism called The Prosperous Society. I'd love for you to walk me through the arguments here. Part one is The Primacy of Distribution, this certainly sounds right in my wheelhouse, so give me the overall view but then I'll ask you about each of the four parts.

ES: Let me give the big picture thesis, which I kind of encapsulate in the fourth episode. Let me give a little bit of the background.

I was writing this as a book, and then the Citrini short report dropped and I just couldn't bite my tongue. So I scrapped the book and I tried to break it up into four podcast parts, the whole running time is like three-and-a-half hours, I'm

going to release one combined episode soon of no ads. I wanted to make it episodic, but I just felt like I just viscerally needed to respond to the Citrini report and so I came out with episode one just as a kind of a response to that and I cobbled it together from what was the manuscript.

My point is when you get the explosion of output that AI is allowing for, you shift the binding constraint to distribution, because you've got this massive pool of stuff to pull from.

You're talking to the right guy, you have an explosion of supply, the problem becomes how do you sort through this? How do you actually find what you want?

ES: Exactly, and so the distribution becomes the principal allocation mechanism, but these allocation mechanisms with digital ads are very, very good, they're very, very efficient. So that's great, you're shifting a lot, a big chunk of the economy to these systems that are very, very efficient and they're very, very workable, and that's good, that automatically creates growth, that automatically creates value.

But the other thing is in that explosion of things that are being created — and it's not just slop output, but it's all types of content. People are getting more efficient at creating great podcasts, great books, great blog posts, and it's not that the AI is writing it, it's just they're getting more efficient at pulling ideas together, doing research, packaging it. You get this larger exploration of the product space, you're creating more and more granular products that could satisfy very specific and particular needs. And then you're getting more products that meet more very specific particular needs, you're getting more efficient matching and you're getting people that are on-boarded to the advertising economy that never could have done so in the first place because of this primacy of distribution that all the ad platforms are investing into making these systems end-to-end automated that brings in ever smaller SMBs [Small and Medium-Sized Businesses] — so across these three axes you're creating more economic value and I think that's what leads to the prosperous society.

Actually the foil for this was this book, I've got it right in front of me, John Galbraith's **Affluent Society** — he wrote this in the 50s, but his point was advertising was just introduced as a means of maximizing consumption because the conventional wisdom at the time, what he called the conventional wisdom, is that we need to be at full economic output all the time, that means producing as much as possible. We're producing more stuff than people demand, and so we need to manufacture demand, and that's what advertising is.

And my point with the prosperous society is actually, the demand-matching system is totally indifferent to what you want, whatever you want, it'll find you the thing that you want most and that you're going to enjoy the most and so people are getting more value and more personal satisfaction out of the things they discover. There's more things to discover, so they're discovering more granular things, and the matching is happening at the maximum possible efficiency. So the advertisers are getting the most value and the customers are getting the most value and the platforms are getting the most value. And in this way, what we can do is we can take that 83% of all commerce, of all retail that's not happening online, and let's say we double that and when we double it, it's much more efficient. All of this distribution is much more efficient, all the consumption is much more efficient and satisfying, that is a prosperous society, that's everyone wins, that's growth across all those axes, and this is not the doomsday scenario that Citrini pointed out, this is actually a much happier, much more economically satisfied society.

It's such a good point. And the last question I had here, I'm going to jump to it right here, there's an aspect where to what extent can you separate your and my relative optimism from your and my relative understanding of and appreciation for the power of individualized advertising? Because there's a bit about this that I think the comparison to Galbraith is exactly right. There is an assumption in there of the power and primacy of the supply side and the stupidity and naivety of the demand side, that it can somehow be induced and can be controlled, and the thesis that I think we have, I'll

speaking for myself, but about the reason why I'm such an advocate for Meta-style advertising is the belief that human desires and needs and wants are effectively infinite and it's infinite in a combinatorial way.

It's not just that humanity broadly wants things, it's that every individual wants different things and the fact that on the production side, we're getting the capability to meet those needs, but we just need to match them, how do we match the need provider to the need haver? And because we believe in advertising, we know it's possible and I think it's not an accident that you have — go back to OpenAI, why are they so resistant to advertising for so long? Because they don't believe it works, they don't understand it and it gets to the whole Silicon Valley blind spot about ads that, by the way, seems to apply to Mark Zuckerberg as well and it falls on us to sit over here and say, “Actually, ads are amazing”. They're not just amazing because they pay for cool toys, they're amazing because they actually understand price and by extension meet consumer needs, and if you understand the massive value explosion that is downstream from ads, it's like, oh my God, AI makes this all a million times better, not worse! It's such a poverty of mindset to think that because AI can just generate a gazillion Soviet-style gray apartment blocks that we've solved the problem, that's not what humans are interested in.

ES: First of all, amen brother. Second, we are, I think, we are like facing this, call it a failure of imagination or call it a moral malaise, but if you look at a lot of the discourse, let's call it and who it's dominated by, it's people that have firmly internalized a degrowth mindset and the affluent society in Galbraith's book is the degrowth manifesto. It's quoted to this day by a lot of these degrowthers and that's really scary to me, and I think it's really important for people to advocate for the opposite.

Capitalism is good, capitalism enriches people, capitalism creates fair societies, capitalism creates happier people, capitalism creates more human satisfaction. And yeah, you could obviously fix things and optimize things at the margin, but

largely it's better, it's objectively better than the alternatives. And my worry here, and if you don't, if someone doesn't actually make the case and actually go through and analytically make the case and make the case with support and make the case by tying it to books that we all agree are Western canon, then if you don't do that, then the degrowth mindset just by default dominates and that's a really scary proposition.

Well, and it's combined with this. People in Silicon Valley are smart, the Communist theorists were always smart, and implicit in that is this sort of looking down on the general consumer and how dumb they are and this lack of imagination for the tremendous capacity of human wants and desires and the willingness and innovative capability of humans to meet those desires. You talked about in part two, The Human Nexus of Commerce and The Collapse of the Pareto Principle, I think both of them touch on this infinite customization aspect and that to the extent we accelerate both sides, that's a huge positive.

ES: And to be fair, in part four, I call out some potential issues with endless atomization and you could think about people becoming alienated and social cohesion breaking down, those are risks, but my bigger worry is not that those take root, because they won't, because consumers will reject that, my bigger worry is that this is what the companies invest all of this CapEx expenditure into, and then it doesn't work.

Then we're in a bubble because we're not in a bubble — if it works, we're not in a bubble, we're on a growth trajectory. If it doesn't work, we're in a bubble and then the bubble bursts. And then we have a recession, and that's a bad outcome. But the thing is, my sense is you've got the lump of labor fallacy, which I think most people recognize that as just broken on its face.

But the lump of labor fallacy is the exact same mindset that leads to this total, "Human labor is going to be worthless", it's like two sides of the same coin.

ES: And then that's coupled with this deep cynicism and nihilism that is really destructive and corrosive and I think you've got to match that with optimism and you have to shout it from the rooftops and I'm doing my part, I know you're doing your part, but the thing is like, if a Citrini publishes this thing and it tanked the S&P the next day, you need to come back with something else, there needs to be a rebuttal, there needs to be a very pointed rebuttal there.

That's essentially what I was trying to do, I hope some people enjoyed it, I just think it's incumbent on the people that have any sort of voice to make this case, because if you don't make the case, the other side's going to dominate.

I like this thesis. People who like ads are optimistic about AI.

ES: There you go.

I think there's honestly something to that. There's this suspicion of commerce, this sort of assumption that people are dumb, that is very much tied in this lack of imagination for the expanse of potential products and services to be sold and acquired that are absolutely absurd. I was on another podcast last week and someone was like, "They made this leap from planting seeds in the ground and eating your own food to losing value in your office job from AI" — I'm like, "Wait, there's a lot of steps here from the intrinsic value of feeding yourself to email jobs", we're far departed from that, and our ability to invent completely new needs and functions and to create solutions to that and then to match them together — it already works pretty well, that's why so many people have missed out on all this Meta upside all these years.

ES: That's right, and Google.

Eric, good talking to you, congratulations on your thesis and very excited to see what happens next.

ES: Always a pleasure Ben, thank you for having me.

This Daily Update Interview is also available as a podcast. To receive it in your podcast player, [visit Stratechery](#).

The Daily Update is intended for a single recipient, but occasional forwarding is totally fine! If you would like to order multiple subscriptions for your team with a group discount (minimum 5), please contact me directly.

Thanks for being a supporter, and have a great day!

Related

An Interview with Eric Seufert about Apple, Facebook, and Mobile Advertising

Thursday, February 4, 2021

2026.22: Luceing Their Mind

Friday, May 29, 2026

2025.47: Gemini! At The Disco

Friday, November 21, 2025

← [The SpaceX IPO and Data Centers in Space](#)

[2026.22: Luceing Their Mind](#) →

Stratechery *Plus* UPDATES

An Interview with Figma CEO Dylan Field About Design and AI

Thursday, June 25, 2026

My Vibe Coding Adventure, The App and the Experience, Ten Takeaways

Wednesday, June 24, 2026

Memory Chips and China, Microsoft and Chinese Models

Tuesday, June 23, 2026

[View All](#)

Stratechery *Plus* PODCASTS


10 Questions on LeBron, Jokic, KESSLER, Blazers Points Guards and Offseason Chaos

Greatest Of All Talk | Jul 1



Price Hikes

Dithering | Jun 26



(Preview) A Summer Break Mailbag: Memory Mania, Vibe Coding, Mafia PR, Caffeine



Intake, Garages, and How to
Fix Soccer
Sharptech | Jun 26

Stratechery *Plus* INTERVIEWS

An Interview with Figma CEO Dylan Field About
Design and AI
Thursday, June 25, 2026

An Interview with Michael Morton About E-
Commerce in the Age of AI
Thursday, June 18, 2026

An Interview with Ben Bajarin About Apple, AI, and
Compute
Thursday, June 11, 2026

[View All](#)

© **Stratechery LLC** 2026 | [Terms of Service](#) | [Privacy Policy](#)

Designed with WordPress. Hosted by Pressable.