

— GROWTH CARTOGRAPHY · PERFORMANCE THINKING · ORG DESIGN

The Hourglass bet.

Org design for the agentic era: software now acts on its own. The org chart that survives is the one that protects human judgment while agents do the work.

— SOFTWARE STOPPED BEING A TOOL AND STARTED BEING A PARTICIPANT

~\$0

the marginal cost of generating content. The assembly-line MarCom department loses its reason to exist

\$15T

in B2B purchases could be agent-mediated by 2028, a Gartner estimate cited in Ramp's analysis of agentic commerce

— PART ONE • PEOPLE

The expert generalist

Specialists owned lanes. Generalists own outcomes, with agents as on-demand domain experts. Hire for first principles, system curiosity, taste, and orchestration.




— THE DESKILLING TRAP

Four shapes an org can take. One survives 2034.

01


The Pyramid



Many juniors, few seniors. Too slow and too heavy for agentic workflows.

02


The Diamond



Juniors cut, middle bloated to watch the AI. Destroys the learning pipeline.

03

The Inverted Pyramid



Senior pods only. Fast today, no path for anyone to become senior tomorrow.

04

The Hourglass



Senior pods on top, lean middle, protected junior base building judgment below.

“

Expert judgment is not innate. It is built by a decade of friction, mistakes, and edge cases. **If nobody junior is in the system today, nobody senior exists in 2034.** The junior layer's output is not the asset. It is the judgment.

— PART TWO · PROCESS

Pods on a paved road

Ops handoffs break on non-deterministic agents. Embedded pods scale badly alone. The target: autonomous pods on a shared platform for runtime, identity, and observability.



— LOOP ENGINEERING · AFTER LANGCHAIN AND ADDY OSMANI

Four loop levels: from task runner to self-improving harness

LEVEL 1

Agent Loop

One model calls tools until a discrete task completes.
Automates isolated, high-volume work.

LEVEL 2

Verification Loop

Output scored against a strict rubric; failures retry with feedback. Brand safety, compliance, factual checks.

LEVEL 3

Event-Driven Loop

External signals trigger runs.
Responsive market action without a human initiating.

LEVEL 4

Hill-Climbing Loop

The system studies its own production traces and improves prompts, models, and memory over time.

— GOVERNANCE AS INFRASTRUCTURE

The river needs a riverbank. Prompts are not controls.

CONTAIN

Blast radius at the OS level

Scope agent permissions at the container or operating system layer, not the runtime. What an agent cannot reach, it cannot break.

COMPILE

Policy as code

Brand rules, legal limits, and budgets become executable boolean conditions (Open Policy Agent is the common engine). Violations block instantly, no ticket queue.

RECORD

Log intent, not just action

Forensics in a non-deterministic world means knowing what the agent tried and why, including what the policy layer refused.

— PART THREE · PRODUCTS

Proof beats promise

The 80/20 vision-to-substance ratio inverted.
The Proof Stack (after a16z crypto): deployed
partnerships, hard data, organic pull, and
validation nobody paid for.



— VISIBILITY · THE READER IS INCREASINGLY A MODEL

Pull became push: SEO to GEO / LEO

TRADITIONAL SEO

GEO / LEO

Real-time crawl and index

→ Static training data plus RAG grounding

Backlinks, keywords, domain rating

→ Citation authority, entity depth, reputable third-party mentions

SERP position and traffic volume

→ AI citation frequency and share of voice in answers

Ten blue links compete

→ LLMs cite 2 to 7 domains per query: a zero-sum fight

— A CMO'S FIRST 90 DAYS

Assess, pilot, scale. Protect the base the whole way.

DAYS 1-30

Assess and contain

Map your org's real shape. Freeze middle-heavy hiring. Pull two or three latent polymaths into a pod. Stand up policy as code before agents touch production.

DAYS 31-60

Pilot loops and proof

Ship the Verification Loop first, then an SEO/GEO loop with automated rollback. Re-sequence every pending campaign around the Proof Stack.

DAYS 61-90

Scale pods, build the base

Move to pods plus platform. Retrain teams on citation authority. Launch the sandboxed junior program. Leaders keep maker time on the calendar.



Structure as code, not slides.

MarCom OS is this framework as a working kit: the Hourglass
org blueprint, the capability calculator, and the riverbank
governance templates. SkillFoundry runs the enforcement. Top
Call grades the signal.

falsedawn.industries/marcom-kit

NAMED SOURCES

AWS EVENTS

ANTHROPIC · F. FUNG

LANGCHAIN

ADDY OSMANI

A16Z CRYPTO

FINANCIAL TIMES

PROFOUND

RAMP